

Clue in Prolog

– A Didactic Example –

Luciano Santos

August 2016

1 Introduction

This is the documentation for a simple script in SWI-Prolog that plays the game Clue¹. This implementation follows a didactic approach, not aimed at creating an advanced AI system that employs complex strategies and human behaviour models to master the game. It simply illustrates how a declarative language can be used to play a relatively simple game based on a certain set of rules.

The implementation is based on the rules for the 2002 version of the game (see PDF on the root folder) and the board on Figure 1.



Figure 1: The game board.

The following principles were observed in this implementation to make it simple:

- no long-term planning – for each action and information received, the agent updates its 'knowledge base' and, on each turn, it makes an independent decision based on the current knowledge, instead of following a planned route;
- no lucky guesses – the agent only makes an accusation if it's certain that it's true;
- no poker face – the agent only acts to acquire more information, and will not make a move or guess for the sole purpose of misleading other players;

¹<http://www.hasbro.com/en-us/toys-games/hasbro-games:clue> (accessed on August, 2016)

- no mind reading – the agent will not infer information from other players actions, except facts that can be logically proven; it will not try to predict how people would or should behave, however, it will assume that everyone will play strategically, e.g., if the player to the left has already shown a certain card before, that card will not be used on a next guess, because a smart player would keep showing the same card over and over again, even if she had a different one to show.

The sections below describe the rationale and the details of this implementation. Section 2 shows how the game board and the current position of each player is stored internally and how the agent finds the shortest path to a given goal. Section 3 describes how the knowledge acquired as the match progresses is represented, and how the game decides the action to take in each turn. Finally, Section 4 explains the predicates that allow the final user to initialize and subsequently interact with the agent.

2 Moving on the Board

The game board is seen internally as a grid of size 24×25 . As illustrated by Figure 2, coordinates are relative to the lower left corner, and start on 0.

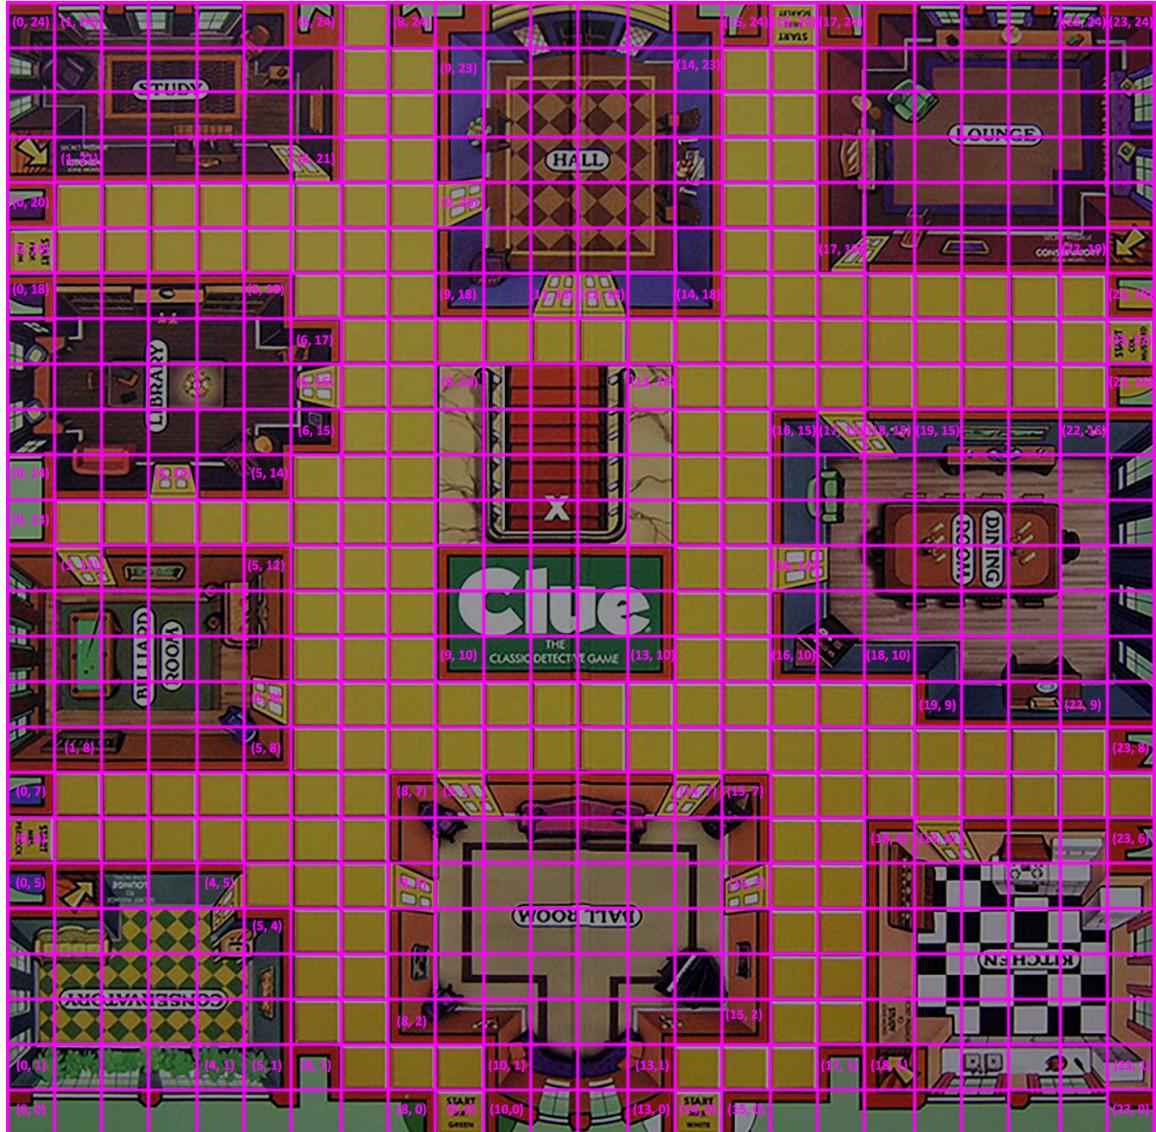


Figure 2: The game board.

2.1 Restrictions

The following predicates describe the restrictions that affect movements inside the board:

- `blocked/1` – is true for the tuple (X, Y) if it represents a location that is blocked, i.e., it's inside a room or in a wall;
- `door/2` – is true for the tuple (X, Y) and room R if it's possible to enter R from (X, Y) ;
- `passage/2` – is true for arguments $\text{Source}, \text{Dest}$ if there is a secret passage from room Source to room Dest ;
- `position/2 (dynamic)` – is true for the tuple (X, Y) and character C if (X, Y) is the current position of C on the board;

Based on the game rules, once inside the room, the position of the individual cell occupied by the player is irrelevant, it's as if each room is a "supercell". For that reason, it's only necessary to know if the player is inside a room and, if not, in which (free) cell she is. Also, because each room is accessible only from specific cells, and each of these cells gives access to one, and only one room, it's unnecessary to know if a given cell belongs to a specific room or wall, it's only important to know which are the access cells (i.e., doors) to which room, and if a given cell is free or not.

Since the blocked areas are more "regular", i.e., they are more easily described in terms of rectangles, it was a design choice to use a predicate `blocked` that defines if a certain position is blocked or not. The same end could be achieved with an opposite predicate `free`. Figure 3 shows the (static) declaration of the blocked cells in the board.

The predicate `door` says if a cell is an access cell, and to which room it gives access to. Since entering the room doesn't count as a step, it's enough to reach these access cells when finding a path and then switching the state to "inside the room". Figure 4 shows the (static) declaration of all the doors in the game.

The predicate `passage` is a special case, to represent the information that there are secret passages between the rooms in the (opposite) corners of the board. As shown in Figure 5, it is implemented as a commutative operation by using an auxiliary predicate.

Finally, the predicate `position` informs the current position of a player. For any atom n that is a valid player name, there will be exactly one fact that says "the position of n is (X, Y) ". The agent makes sure this holds true by defining the start position shown on the board once (Figure 6), and *retracting* and *asserting* that position any time a player's position changes, as will be explained on Section 3. When a player enters a room, her position becomes that room's name.

Other restrictions, such as the board boundaries, are checked on the path finding algorithm, as described on the next section.

2.2 Path Finding

One of the key actions the agent must perform is movement. The first step necessary to move on the board is to answer the question: "where can I go from my current position?". If the player is in a cell adjacent to a door, she can enter the room; if she is in a corner room, she could use a secret passage. However, to get to the point where she could enter a room, the player must first reach a door, and using a secret passage is a trivial action. Thus, right now, only the problem of moving from a free cell to another free cell will be addressed. In order to solve that problem, the agent must determine the shortest path between two points.

As the game rules state, if the player is currently in any given cell, she can only move to another empty cell vertically or horizontally. To express that relationship, the predicate `neighbor/2` is defined (Figure 7). It will be true for points (X_s, Y_s) , (X_n, Y_n) if (X_n, Y_n) is exactly one cell away from (X_s, Y_s) , moving either vertically or horizontally (but not both).

```
%% neighbor((Xs, Ys), (Xn, Yn)) - <Xn, Yn> is neighbor of <Xs, Ys>.
neighbor((Xs, Ys), (Right, Ys)) :- Right is Xs + 1.
neighbor((Xs, Ys), (Left, Ys)) :- Left is Xs - 1.
neighbor((Xs, Ys), (Xs, Up)) :- Up is Ys + 1.
neighbor((Xs, Ys), (Xs, Down)) :- Down is Ys - 1.
```

Figure 7: The declaration of the predicate `neighbor`.

Now, besides checking if a certain cell is a neighbor, it must also be possible to actually move there, *i.e.*, it must be inside the boundaries of the board and not occupied by any other player or blocked, as previously defined. All these cases are summarized by the predicate `is_free/1`, shown in Figure 8.

```
%% is_free((X, Y)) - the position <X, Y> can bee occupied by a character.
is_free((X, Y)) :-
    X >= 0, X <= 23, Y >= 0, Y <= 24, % inside the board
    \+ position((X, Y), _), % not currently occupied by anyone
    \+ blocked((X, Y)). % not a room or a wall
```

Figure 8: The declaration of the predicate `is_free`.

Once all the restrictions on movement between two adjacent cells have been represented, it's possible to implement the algorithm to find the shortest path between two points on the map, if such a path exists. More precisely, the algorithm here implemented finds the shortest path between a given free cell (X_s, Y_s) and the closest door. Subsequent backtracking on the predicate will return all the doors of the board, ascendingly ordered by path length.

The simplest way to tackle this task is using BFS (Breadth-first Search). The idea is simple: inspect the source node; then all the nodes exactly 1 step away from the source node; then all the nodes exactly 2 steps away from the source node; and so on. This ensures that, if a path exists, it will be found and will be a shortest path. If the backtracking mechanism continues from the state where it found a path, the next door found will have the next shortest path.

To make sure that the nodes are inspected in the correct order, a queue is used. This queue begins with the source node. At each iteration, the node at the head of the queue is removed and inspected. If it's not a door, then its adjacent nodes are enqueued and the search continues. To prevent the search from recursing infinitely, it's necessary to keep a set of all nodes seen so far, so only adjacent nodes not yet inspected are enqueued.

The code for this search algorithm is shown in Figure 9. The predicate `closest_door/3` receives a source point and unifies the name of the room with the closest door and the path to that door as a list of points. This predicate itself only sets up the initial values for the algorithm – the queue and the “seen nodes” set are both initialized to a unitary list containing the source point – and calls the auxiliary predicate `closes_door_aux/4` to actually perform the search.

The base case for this predicate is when the head of the queue is a door. If that happens, it will unify to the found door's room and the path is a unitary list containing just the door. The recursive step will extract the head of the queue, enqueue all adjacent cells that are free

and have not yet been seen (described by `valid_adjacent/3`), mark all the neighbor cells are seen, and recursively continue walking on the list. As the recursive calls return, the path is built by inserting the inspected node on the head of the result list.

```

%% closest_door((Xs, Ys), Room, Path) - starting at <Xs, Ys>, finds
%% the closest Room door and the Path to it
unseen_neighbor((X, Y), (Xd, Yd), Seen) :-
    neighbor((X, Y), (Xd, Yd)),
    \+ member((Xd, Yd), Seen).
valid_adjacent((X, Y), (Xd, Yd), Seen) :-
    unseen_neighbor((X, Y), (Xd, Yd), Seen),
    is_free((Xd, Yd)).
closest_door_aux(Room, [Head], [Head|_], _) :- door(Head, Room).
closest_door_aux(Room, [(X, Y)|PTail], [(X, Y)|QTail], Seen) :-
    % enqueues all non-seen adjacents to which it's possible to move
    findall((Xd, Yd), valid_adjacent((X, Y), (Xd, Yd), Seen), ValidAdjacents),
    append(QTail, ValidAdjacents, NewQueue),
    % stores all adjacent nodes as seen, to cut the search recursion
    findall((Xa, Ya), unseen_neighbor((X, Y), (Xa, Ya), Seen), Neighbors),
    append(Neighbors, Seen, NewSeen),
    closest_door_aux(Room, PTail, NewQueue, NewSeen). % recursive definition
closest_door((Xs, Ys), Room, Path) :-
    is_free((Xs, Ys)),
    closest_door_aux(Room, Path, [(Xs, Ys)], [(Xs, Ys)]).

```

Figure 9: The breadth-first search used to find the closest door, given a source cell.

Even though this algorithm finds the path between two cells on the board, the agent actually needs to solve the more general problem of finding the closest doors when departing from multiple source cells. That's because, if the agent is inside a room, there could be multiple exits to choose from, thus, it should find the best option taking into account all these exits.

The naïve approach would be to find lists with all the paths for each exit, merge these lists while ordering the elements (paths) by length and, finally, traverse the result. However, that solution is both inefficient and *imperative*, in the sense that it tells the steps to solve the problem, instead of taking advantage of the declarative paradigm.

What happens if the queue in the previous solution is initialized with multiple source cells, instead of just one?

- when the first element of the queue is inspected, all its adjacent nodes will be enqueued **after** all the other sources, which are currently on the queue (as long as those adjacent nodes aren't sources themselves);
- when the next element is inspected (another source), once again, all its *not already on the queue* adjacents will also be enqueued;
- after all the sources nodes are inspected first, the next element to be inspected is exactly one step away from one of the sources; actually, all the elements that are exactly one step away from at least one of the sources will be inspected next, and their adjacent nodes will be enqueued;
- by analogy, all the elements that are exactly two steps away from **at least one** of the sources will be analyzed next, and so on...

Ergo, it's possible to conclude that, just by initializing the queue with the multiple sources, the order in which the elements are inspected still ensures that doors are searched *breadth-first* and, as a consequence, the doors are found in order, from the closest to the farthest, only now the distance applies to any one of the possible sources.

There's one problem though: in the previous algorithm, since there was only one source, the path could be rebuilt implicitly within the backtracking mechanism. The base of the recursion was the door, and the recursive steps would attach one cell to the head of the resulting list, until the caller of the predicate would have the whole path ready. With multiple sources, however, it's necessary to explicitly store, for each cell that is reached, that cell's parent, *i.e.*, the cell *from which* it was reached.

Figure 10 shows the new version of the algorithm that includes these changes. Now, the Seen set elements are tuples (Cell, Parent) that store not only the coordinate of the cells that were already seen, but also their parent (which can be the coordinates of a previously seen cell or `nil` for the source cells). All predicates were updated to handle this new representation.

The auxiliary predicate `add_parent/3` simply receives a list Nodes and an element Parent and generates a new list with each element N of Nodes in a tuple (N, Parent).

Finally, the new recursive predicate `build_path/3` is used to explicitly rebuild the path when a door is reached, given the set of seen nodes (and their parents). The predicate inspects the Head element using the Seen set and generates the Path as a list; the base case happens when the parent of Head is `nil`, which generates the unitary list; the recursive case adds Head to the list and inspects its non-nil parent. Since this predicate is called when a door is found, the path is actually generated in the reverse order, so the resulting list is reversed before being returned to the original caller.

```

%% closest_door((Xs, Ys), Room, Path) - starting at <Xs, Ys>, finds
%% the closest Room door and the Path to it
unseen_neighbor((X, Y), (Xd, Yd), Seen) :-
    neighbor((X, Y), (Xd, Yd)),
    \+ member(((Xd, Yd), _), Seen).
valid_adjacent((X, Y), (Xd, Yd), Seen) :-
    unseen_neighbor((X, Y), (Xd, Yd), Seen),
    is_free((Xd, Yd)).
add_parent([], _, []).
add_parent([Node|NodesTail], Parent, [(Node, Parent)|LinkedTail]) :-
    add_parent(NodesTail, Parent, LinkedTail).
build_path(Head, Seen, [Head]) :- member((Head, nil), Seen).
build_path(Head, Seen, [Head|Tail]) :-
    member((Head, Parent), Seen),
    build_path(Parent, Seen, Tail).
closest_door_aux(Room, Path, [Head|_], Seen) :-
    door(Head, Room),
    build_path(Head, Seen, InversePath),
    reverse(InversePath, Path),
    print(Path), nl.
closest_door_aux(Room, Path, [(X, Y)|QTail], Seen) :-
    % enqueues all non-seen adjacents to which it's possible to move
    findall((Xd, Yd), valid_adjacent((X, Y), (Xd, Yd), Seen), ValidAdjacents),
    append(QTail, ValidAdjacents, NewQueue),
    % stores all adjacent nodes as seen, to cut the search recursion
    findall((Xa, Ya), unseen_neighbor((X, Y), (Xa, Ya), Seen), Neighbors),
    add_parent(Neighbors, (X, Y), LinkedNeighbors),
    append(LinkedNeighbors, Seen, NewSeen),
    closest_door_aux(Room, Path, NewSeen, NewSeen). % recursive definition
closest_door((Xs, Ys), Room, Path) :-
    is_free((Xs, Ys)),
    closest_door_aux(Room, Path, [(Xs, Ys)], [(Xs, Ys), nil]). % recursive definition
closest_door(SourceRoom, TargetRoom, Path) :-
    valid_rooms(ValidRooms), member(SourceRoom, ValidRooms),
    findall(Exit, door(Exit, SourceRoom), Exits),
    add_parent(Exits, nil, LinkedExits),
    closest_door_aux(TargetRoom, Path, Exits, LinkedExits).

```

Figure 10: The breadth-first search used to find the closest door, given multiple source cells.

3 Making Decisions

This section describes how the agent represents knowledge about the game internally (Section 3.1), and how it then uses that knowledge to make decisions in each turn (Section 3.2).

3.1 Representing the Game Data

The first necessary data to represent internally are the lists of valid rooms, characters and weapons. These (static) lists are initialized with the predicates in Figure 11 and used throughout the code. For simplicity, no (physical) player name is stored, instead, all players are referenced by their character's name. The list of valid characters is declared in the order that those characters would play, according to the rules. The agent's own character is represented by the dynamic predicate `my_char/1`, that is set once when a match starts.

```
%% valid_rooms(Rooms) - Rooms is the list of valid room names.
valid_rooms(
[

    'conservatory',
    'ball room',
    'kitchen',
    'dining room',
    'lounge',
    'hall',
    'study',
    'library',
    'billiard room'
].
).

%% valid_chars(Chars) - Chars is the list of valid characters,
%% in the order they must play.
valid_chars(
[

    'scarlet',
    'mustard',
    'white',
    'green',
    'peacock',
    'plum'
].
).

%% valid_weapons(Weapons) - Weapons is the list of valid weapons.
valid_weapons(
[

    'rope',
    'pipe',
    'knife',
    'wrench',
    'candlestick',
    'pistol'
].
).

%% my_char(Char) - Char is the character of this agent.
:- dynamic my_char/1.
```

Figure 11: The lists of rooms, characters and weapons in the game; and the current player's char.

Next, an important information required by the agent to make a move is “which cards have I seen?”. To win the game, a player uses a very simple elimination process: the total number of cards is known for each category – room, character and weapon; there's exactly one card of each in the “confidential” envelope; if, for any of those, only a single card has not yet been seen, then it's logical that it must be on the envelope.

Figure 12 shows the very straightforward predicates that represent the information of which cards have been shown, and by which player. This last information is necessary for the decision step described later. Notice that the player who has shown the card might be the agent itself, meaning that the card is on its hands.

The predicate `can_accuse/3` is true if the agent currently knows all the necessary information to make an accusation. It's trivially defined as the case when the unknown rooms, characters and weapons lists are all unitary.

```

%% shown_char(Player, Char) - the player Player (could be myself)
%% has shown me the card for Char.
:- dynamic shown_char/2.

%% shown_room(Player, Room) - the player Player (could be myself)
%% has shown me the card for Room.
:- dynamic shown_room/2.

%% shown_weapon(Player, Weapon) - the player Player (could be myself)
%% has shown me the card for Weapon.
:- dynamic shown_weapon/2.

%% the next predicates generate the known and unknown lists
%% of characters, rooms and weapons, i.e., those whose cards
%% have and have not yet been shown so far
known_chars(KnownChars) :- findall(C, shown_char(_, C), KnownChars).
unknown_chars(UnknownChars) :-
    valid_chars(ValidChars),
    known_chars(KnownChars),
    subtract(ValidChars, KnownChars, UnknownChars).

known_rooms(KnownRooms) :- findall(C, shown_room(_, C), KnownRooms).
unknown_rooms(UnknownRooms) :-
    valid_rooms(ValidRooms),
    known_rooms(KnownRooms),
    subtract(ValidRooms, KnownRooms, UnknownRooms).

known_weapons(KnownWeapons) :- findall(C, shown_weapon(_, C), KnownWeapons).
unknown_weapons(UnknownWeapons) :-
    valid_weapons(ValidWeapons),
    known_weapons(KnownWeapons),
    subtract(ValidWeapons, KnownWeapons, UnknownWeapons).

can_accuse(Person, Room, Weapon) :-
    unknown_chars([Person]),
    unknown_rooms([Room]),
    unknown_weapons([Weapon]).
```

Figure 12: Representing which cards were shown, by whom.

3.2 Picking Actions

4 The Interface

```

%% blocked((X, Y)) - point <X, Y> is inside a room or is a wall
% conservatory
blocked((X, Y)) :- Y == 0, X >= 0, X <= 8.
blocked((X, Y)) :- X >= 0, X <= 4, Y >= 1, Y <= 5.
blocked((X, Y)) :- X == 5, Y >= 1, Y <= 4.
blocked((6, 1)).
% ball room
blocked((X, Y)) :- X >= 10, X <= 13, Y >= 0, Y <= 1.
blocked((X, Y)) :- X >= 8, X <= 15, Y >= 2, Y <= 7.
% kitchen
blocked((X, Y)) :- Y == 0, X >= 15, X <= 23.
blocked((17, 1)).
blocked((X, Y)) :- X >= 18, X <= 23, Y >= 1, Y <= 6.
% dining room
blocked((X, Y)) :- X >= 16, X <= 18, Y >= 10, Y <= 15.
blocked((X, Y)) :- X >= 19, X <= 22, Y >= 9, Y <= 15.
blocked((X, Y)) :- X == 23, Y >= 8, Y <= 16.
% lounge
blocked((X, Y)) :- X >= 17, X <= 22, Y >= 19, Y <= 24.
blocked((X, Y)) :- X == 23, Y >= 18, Y <= 24.
% hall
blocked((X, Y)) :- X >= 9, X <= 14, Y >= 18, Y <= 23.
blocked((X, Y)) :- Y == 24, X >= 8, X <= 15.
% study
blocked((X, Y)) :- X == 0, Y >= 20, Y <= 24.
blocked((X, Y)) :- X >= 1, X <= 6, Y >= 21, Y <= 24.
% library
blocked((X, Y)) :- X >= 0, X <= 5, Y >= 14, Y <= 18.
blocked((X, Y)) :- X == 6, Y >= 15, Y <= 17.
% billiard room
blocked((X, Y)) :- X == 0, Y >= 7, Y <= 13.
blocked((X, Y)) :- X >= 1, X <= 5, Y >= 8, Y <= 12.
% stairs
blocked((X, Y)) :- X >= 9, X <= 13, Y >= 10, Y <= 16.

```

Figure 3: The declaration of the predicate `blocked`.

```

%% door((X, Y), R) - there's a door to room R from point <X, Y>
door((5, 5), 'conservatory').
door((7, 5), 'ball room').
door((16, 5), 'ball room').
door((9, 8), 'ball room').
door((14, 8), 'ball room').
door((19, 7), 'kitchen').
door((15, 12), 'dining room').
door((17, 16), 'dining room').
door((17, 18), 'lounge').
door((8, 20), 'hall').
door((11, 17), 'hall').
door((12, 17), 'hall').
door((6, 20), 'study').
door((3, 13), 'library').
door((7, 16), 'library').
door((1, 13), 'billiard room').
door((6, 9), 'billiard room').

```

Figure 4: The declaration of the predicate `door`.

```

%% passage(Src, Dst) - there's a passage from Src to Dst
passage_aux('conservatory', 'lounge').
passage_aux('kitchen', 'study').
passage(Src, Dst) :- passage_aux(Src, Dst) ; passage_aux(Dst, Src).

```

Figure 5: The declaration of the predicate `passage`.

```

%% position((X, Y), C) - the current position of character C is <X, Y>
%% here, it's initialized to the start position at the board.
%% this predicate will be "rewritten" whenever the script's own
%% character moves or it receives information that another character
%% moved.
:- dynamic position/2.
position((16, 24), 'scarlet').
position((23, 17), 'mustard').
position((14, 0), 'white').
position((9, 0), 'green').
position((0, 6), 'peacock').
position((0, 19), 'plum').

```

Figure 6: The declaration of the predicate `position` and the initial position of all players.